

Sawyer Czupka

+1 (757) 818-0742 — sawyerczupka@gmail.com — sawyerczupka.com — linkedin.com/in/sawyerczupka

Recent CS & Data Science graduate with 2+ years of experience spanning full-stack engineering, cloud infrastructure, and machine learning—from training and evaluating models to deploying them at scale—with a track record of owning complete technical stacks end-to-end.

SKILLS

Programming & ML/Modeling

Python, TypeScript, JavaScript, Java, C++, Rust, PyTorch, PyTorch Geometric, Scikit-Learn, Pandas, Hugging Face Transformers, React, Vue, FastAPI, Django, SQLAlchemy, PostgreSQL, TimescaleDB, MongoDB, Qdrant, LlamaIndex, DeepEval, vLLM, Azure OpenAI, Git

DevOps & Infrastructure

Docker, Kubernetes, Terraform, Packer, ArgoCD, GitHub Actions, CI/CD, Prometheus, Grafana, Traefik, MLflow, Azure, AWS, AWS Lambda, Linux, Nginx, Let's Encrypt

EDUCATION

College of William & Mary

BS in Computer Science & Data Science, Cum Laude & Dean's List

Sept. 2021 – May 2025

GPA: 3.6

EXPERIENCE

Software Engineering Intern (Sole Engineer), Luminexis AI / Threat Tec

Aug. 2025 – Dec. 2025

- Served as sole engineer owning the entire stack; built production React/FastAPI platform from scratch with JWT auth, role-based access control, admin dashboards, and ROI tracking tools for AI automation opportunity identification.
- Engineered LLM-powered document analysis pipeline using Azure OpenAI and Document Intelligence to extract requirements, risk factors, and action items from 140+ contract PDFs, reducing manual review from hours to minutes.
- Deployed RAG chatbot (Qdrant, LlamaIndex) achieving 70–80% on RAGAS evaluation metrics; provisioned 4-environment Azure infrastructure with Terraform and fully automated CI/CD via GitHub Actions.

Machine Learning & Software Engineer, Teamculture.ai / L10.tech

Jan. 2024 – Sept. 2024

- Built serverless RAG evaluation system on AWS Lambda and API Gateway using DeepEval, eliminating manual AI performance testing through fully automated pipeline execution.
- Designed Vue/FastAPI interface for creating evaluations and curating golden example datasets, enabling systematic quality measurement and continuous improvement of AI services at scale.

Machine Learning Technical Lead, GeoLab @ William & Mary

Jan. 2023 – May 2025

- Led ML integration across frontend, backend, and ML subteams for SCOPE research platform, defining technical architecture and coordinating end-to-end implementation.
- Architected microservice RAG system on Kubernetes (Qdrant, vLLM, LlamaIndex, FastAPI); built Vision-Language Model document processing pipeline (Qwen-2.5-VL) achieving 50%+ reduction in parsing errors vs. conventional tools on complex multi-format documents.

Machine Learning & Software Engineer Intern, The World Bank & GEF

May 2023 – Sept. 2023

- Built RAG system automating analysis of 24,000+ GEF project documents for socioeconomic impact evaluation, enabling data-driven funding allocation and strategic planning at a scale previously infeasible manually.

PERSONAL PROJECTS

LapEvo: iRacing Telemetry Tracker & Real-time Coach

Jan. 2025 – Present

- Architected real-time 60Hz telemetry pipeline with Rust client, FastAPI/TimescaleDB backend, and React 19/TypeScript dashboard; implemented pytest/testcontainers integration tests enabling CI/CD without live game dependencies.

Personal Homelab Infrastructure

July 2023 – Present

- Operate production-grade homelab with 6 VMs and 116-pod Kubernetes cluster on Proxmox, hosting services behind Traefik/Let's Encrypt; manage full IaC stack using Packer, Terraform, Kubernetes manifests, and ArgoCD for GitOps continuous deployment.

ACADEMIC PROJECTS

Selected ML Projects

William & Mary, 2021–2025

- Molecular property prediction with GNNs (PyTorch Geometric); transformer-based sentiment analysis; CNN image classification and regression; ML-based network traffic trace classifier; time series forecasting; generative modeling with diffusion models.